

# On Optimally Partitioning a Text to Improve Its Compression [1]

Davide Cologni<sup>1</sup>, PhD Student

Venice, Italy – 24 January 2026

---

<sup>1</sup>RAVEN Group – Ca' Foscari University of Venice,  
Github: [github.com/colobrodo](https://github.com/colobrodo),  
Email: [davide.cologni@unive.it](mailto:davide.cologni@unive.it)

# Problem: Text Partitioning

We have a **compressor**  $\mathcal{C}$  and a **text**  $T$  of size  $n$ , is it possible to **divide**  $T$  into  $k \leq n$  parts,  $T[1..i_1 - 1]T[i_1..i_2 - 1]\dots T[i_{k-1}..n]$  and **compress each** of them individually with  $\mathcal{C}$  to improve the overall compression?

**Note:** We do **not** *permute* the string.  
We are only interested in *partitioning* it.

# Text Partitioning Example

Suppose we have the text  $T = a^n b^n$ .


If we compress the entire text at once we should use one bit per symbol, or  $O(n)$  **bits**.

If instead we partition the text to compress  $a^n$  and  $b^n$  separately we can compress the whole string using only  $O(\log_2(n))$  **bits** indicating just the length of each substring.

# Reduction to SSSP

We can model the partition problem as a **directed graph** with  $n + 1$  vertices, where an edge exists between  $v_i$  and  $v_j$  only if

$$1 \leq i < j \leq n + 1$$



# Reduction to SSSP - Bijection between paths and partitions

In this graph each **edge** represents a **substring** of the text. We can then show that there exists a **bijection** from each **path**  $\pi = (v_1, v_{i_1}) \dots (v_{i_k}, v_{n+1})$  in the graph, and a **partitioning** of the text  $T$  in the form  $T[1..i_1 - 1]T[i_1..i_2 - 1]\dots T[i_{k-1}..n]$

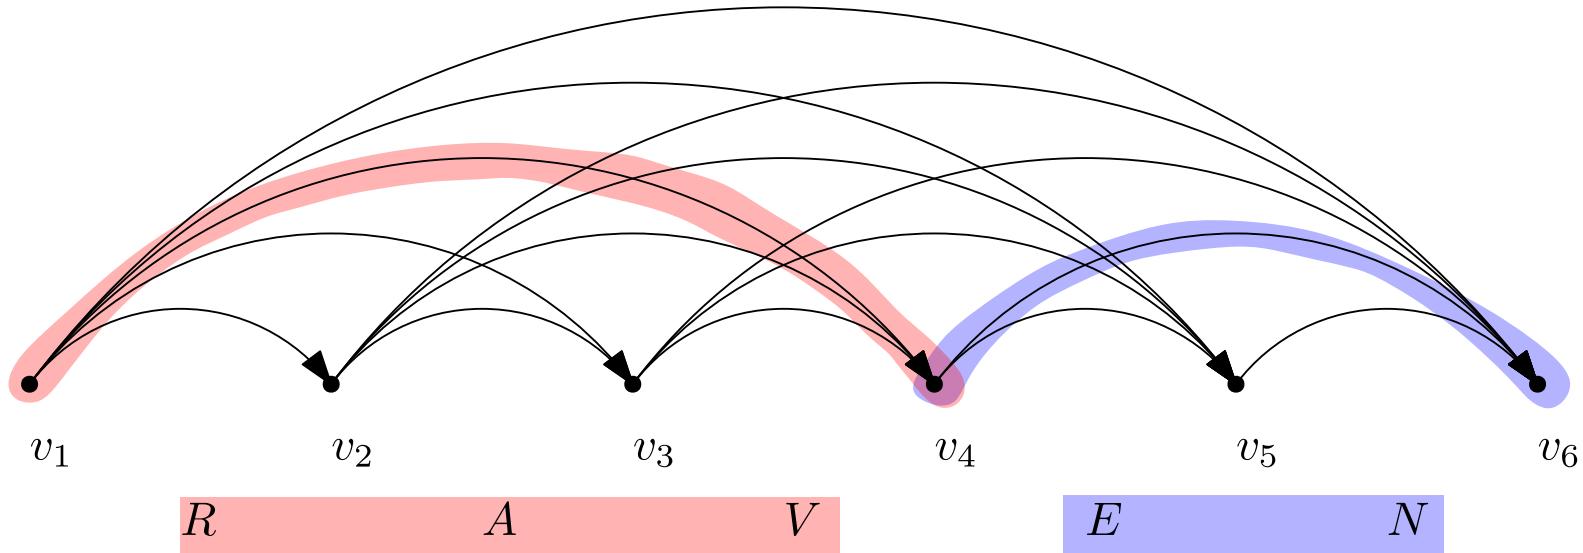


Figure 2: We can map the path  $\pi = (v_1, v_4) | (v_4, v_6)$  to the partitioning of the string  $T[1, 3] | T[4, 5]$

# Reduction to SSSP - Bijection between paths and partitions

If we weight each edge  $(i, j)$  of the graph by the cost of compressing the corresponding text segment  $w(i, j) = |\mathcal{C}(T[i, j - 1])|$ , we can solve the partitioning problem *optimally* computing the **Single Source Shortest Path (SSSP)**

It can be computed efficiently in  $O(|E|)$  time using a classic dynamic programming algorithm.

# Problems:

1. Our graph has  $O(n^2)$  nodes by construction
2. To initialize the weight  $w(i, j)$  we should execute  $\mathcal{C}$  on every substring of the text

# Assumption on $\mathcal{C}$

- Our compressor is *monotonic*: the compressed output on a suffix or a prefix of the string is always smaller than the compression on the whole string:

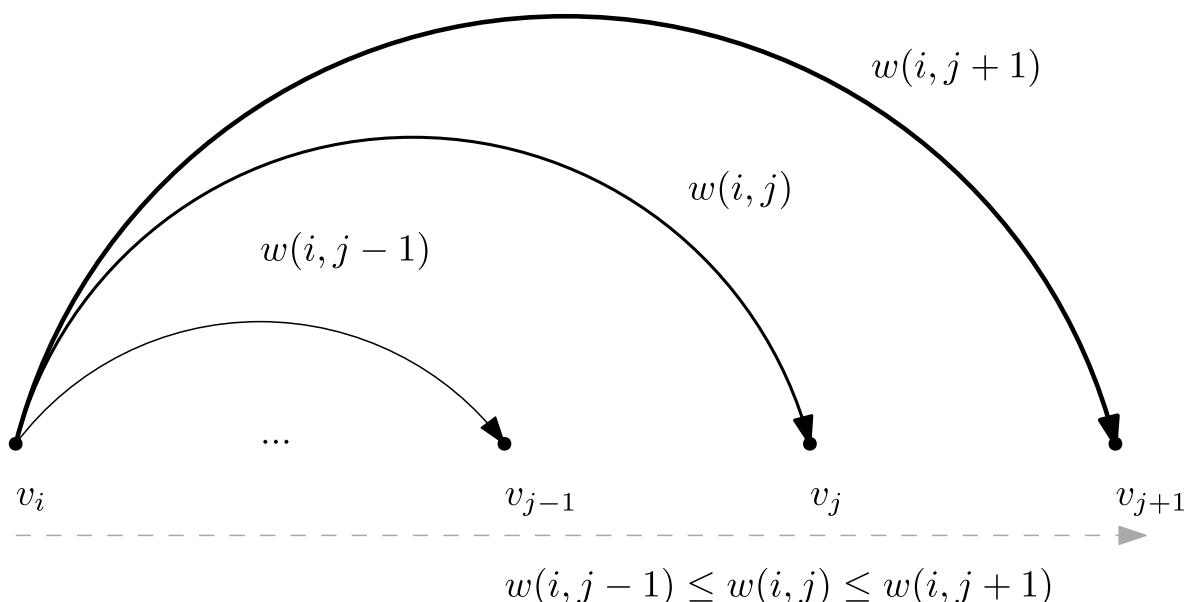
$$|\mathcal{C}(T[i, j])| \geq |\mathcal{C}(T[i, j - 1])|$$

$$|\mathcal{C}(T[i, j])| \geq |\mathcal{C}(T[i + 1, j])|$$

- We can compute the size of the compressed output incrementally: computing  $|\mathcal{C}(T[i, j])|$  from the state of  $\mathcal{C}(T[i - 1, j])$  or  $\mathcal{C}(T[i, j - 1])$  takes constant time

# Monotonicity of $w$

Due to the monotonicity of the compressor for every node  $1 \leq i < k < j \leq n + 1$  we have that  $w(i, k) \leq w(i, j)$



# Sparsification of the DAG

Thanks to this property we can obtain an approximated algorithm by **sparsifying** the graph thus selecting only some edges.

We are able to obtain a  **$(1 + \varepsilon)$ -approximation**, for every  $\varepsilon \geq 0$ , with a time complexity of  $O(n \log_{1+\varepsilon} L)$

where  $L = w(1, n)$ , so the cost of compressing the entire text.

This algorithm can be applied to every dynamic programming algorithm in the form  $E[j] = \min_{1 \leq i < j} (E[i] + w(i, j))$  when  $w$  is *monotone!*


# Key Idea: $\varepsilon$ -maximal edges

How we can select some edges to obtain the  $(1 + \varepsilon)$  approximation factor?

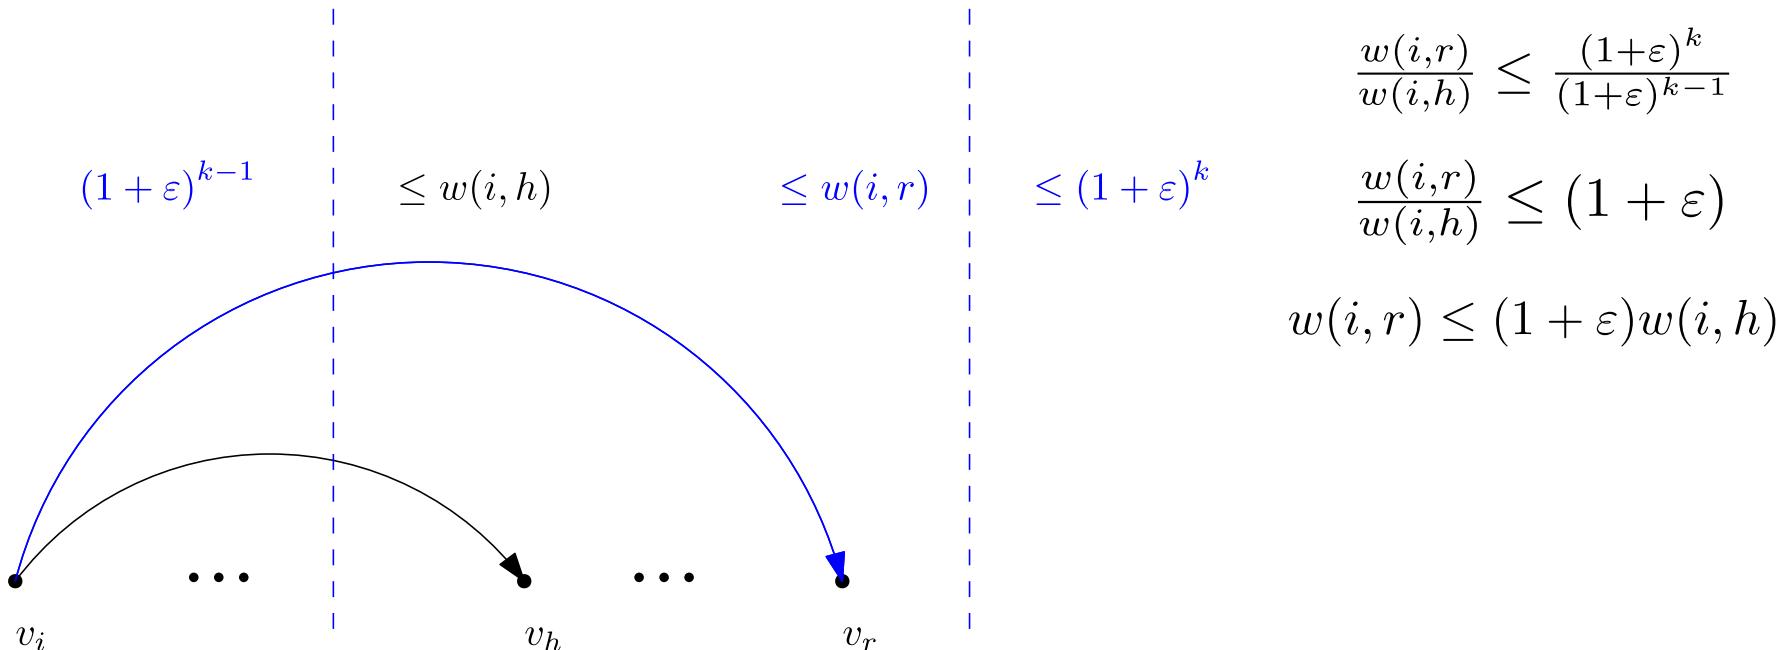
For each node  $i$  select the  $\varepsilon$ -maximal edges, so the outgoing edge from  $i$  that satisfy one of these conditions:

- The edges  $(i, j)$  such that  $w(i, j) \leq (1 + \varepsilon)^k < w(i, j + 1)$  for any integer  $k \geq 1$
- The last outgoing edge:  $(i, n + 1)$

So we select the best approximations of the powers of  $(1 + \varepsilon)$  from below: We then have at most  $\log_{1+\varepsilon} L$  outgoing edges for each node.



Each edge is then “*covered*” by an  $\varepsilon$ -maximal edge: The weight of the edge is then approximated by  $(1 + \varepsilon)$  times the weight of the maximal edge that covers it.




**Our edges are increasing and can be approximated but what can we say about *paths* in this graph?**

## Lemma 1

Let  $d_{\mathcal{G}}(i)$  be the cost of the shortest path  $\pi_i$  in our graph  $\mathcal{G}$  from  $v_i$  to  $v_{n+1}$  then

For all the vertices  $i, j : 1 \leq i < j \leq n + 1$ ,  $d_{\mathcal{G}}(i) \geq d_{\mathcal{G}}(j)$



# Theorem

Let  $\mathcal{G}$  be the full graph and  $\mathcal{G}_\varepsilon$  be the graph containing only  $\varepsilon$ -maximal edges, then  $d_{\mathcal{G}_\varepsilon}(i) \leq (1 + \varepsilon)d_{\mathcal{G}}(i)$  for every integer  $1 \leq i \leq n + 1$ .

## Proof by induction:

- **Base**, trivial case for  $n + 1$
- Then let  $\pi(i) = (v_i, v_{t_1})..(v_{t_h}, v_n)$  the shortest path starting from node  $v_i$  and let  $d_{\mathcal{G}}(i) = w(i, t_1) + d_{\mathcal{G}}(t_1)$  be its cost. We choose the  $\varepsilon$ -maximal node  $r$  that covers  $t_1$ : So  $r > t_1$  and we already know (by our "key idea") that

$$w(i, r) \leq (1 + \varepsilon)w(i, t_1)$$

By *Lemma 1*:


$$d_{\mathcal{G}}(r) \leq d_{\mathcal{G}}(t_1)$$

By inductive hypothesis:

$$d_{\mathcal{G}_\varepsilon}(r) \leq (1 + \varepsilon)d_{\mathcal{G}}(r) \leq (1 + \varepsilon)d_{\mathcal{G}}(t_1)$$

In the end

$$d_{\mathcal{G}_\varepsilon}(i) = w(i, r) + d_{\mathcal{G}_\varepsilon}(r) \leq (1 + \varepsilon)(w(i, t_1) + d_{\mathcal{G}}(t_1)))$$



# Problem: DAG Construction


We still have two problems:

1. if we construct *naively* this graph we should remove edges from a  $O(n^2)$  graph
2. and we should compute the weight of each edge of the graph

*We can solve both these problems efficiently at once:* We can find the  $\varepsilon$ -maximal edges efficiently on the fly!

# Sliding windows

We keep  $\log_{1+\varepsilon} L$  sliding windows all starting at  $v_i$ , but ending in a different position. The  $k$ -th window find the  $k$ -th  $\varepsilon$ -maximal edge.




$w_{k-1}$

$w_k$

$w_{k+1}$

# Sliding windows


For each compressor we should implement 2 operations on the windows `advance_left`, `advance_right`: The first operation advances the start of **all** the windows.



# Sliding windows

`advance_right` advance the end of the  $k$ -th window of one position.

We call this function until we reach the last edge smaller than  $(1 + \varepsilon)^k$ , so until we find the  $k$ -th maximal edge starting from node  $i$ .



if the operations `advance_left` and `advance_right` have respectively a complexity of  $O(L)$  and  $O(R)$  our algorithm execute asymptotically  $O(Ln + Rn \log_{1+\varepsilon} n)$  steps

The authors provide several implementations of the sliding windows framework to estimate the size of different compressors, among the others statistical compressors (using 0-th order and k-order entropy)

# Computing Zero Order Entropy

Zero-th order entropy is a well-known lower bound for the performance of statistical compressors.

For each windows  $w_i$ , we maintain a histogram,  $A_i[c]$ , indexed by the symbol  $c \in \Sigma$

$$E_i = \sum_{c \in \Sigma} A_i[c] \log_2 A_i[c]$$

Using  $E_i$ , we can calculate a lower bound on the output of the statistical compressor,  $|\mathcal{C}(T[..i])|$  based on the zero-th order entropy as

$$|T[..i]| H_0(T[..i]) = |T[..i]| \log_2 |T[..i]| - E_i$$

From this we can calculate incrementally the value of  $E_{i+1}$  removing the old term from the summation and adding the new one:

Let  $c = T[i + 1]$  then

$$E_{i+1} = E_i - A_i[c] \log_2 A_i[c] + (A_i[c] + 1)(\log_2 A_i[c] + 1)$$

# Thank You!

# Bibliography


- [1] P. Ferragina, I. Nitto, and R. Venturini, "On Optimally Partitioning a Text to Improve Its Compression," in *Algorithms - ESA 2009*, A. Fiat and P. Sanders, Eds., Berlin, Heidelberg: Springer Berlin Heidelberg, 2009, pp. 420–431.
- [2] G. Ottaviano and R. Venturini, "Partitioned Elias-Fano Indexes," in *Proceedings of the 37th International ACM SIGIR Conference*, in SIGIR '14. New York, NY, USA: ACM, 2014, pp. 273–282. doi: [10.1145/2600428.2609615](https://doi.org/10.1145/2600428.2609615).

# Bonus Slides: Partitioned Elias-Fano

## Elias-Fano Data Structure

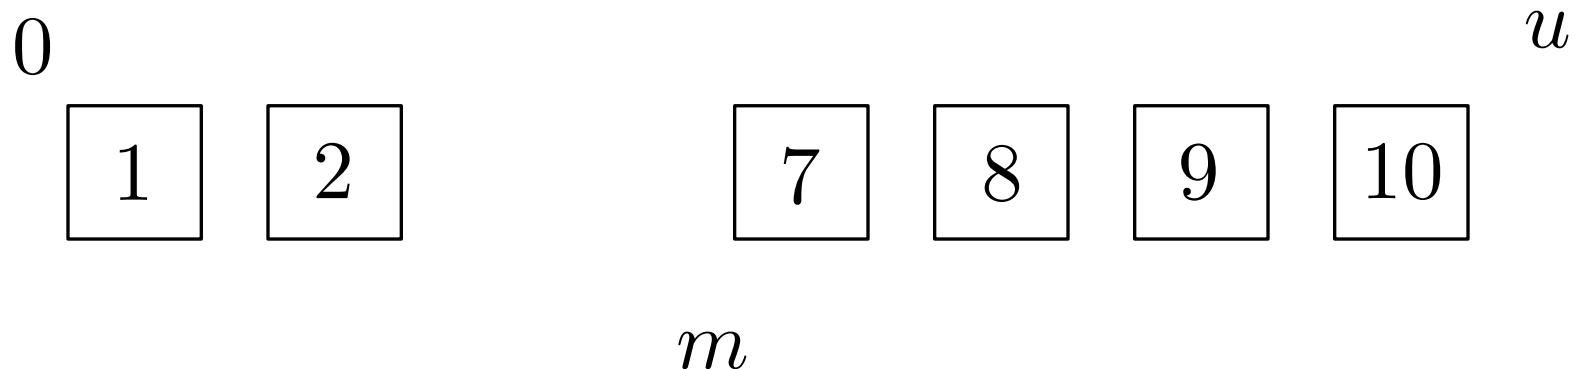
A compact data structure to store a set of  $m$  monotonically increasing integers upper-bounded by  $u$ .

It uses  $\approx \lceil \log_2 \frac{u}{m} \rceil + 2$  bits per element.



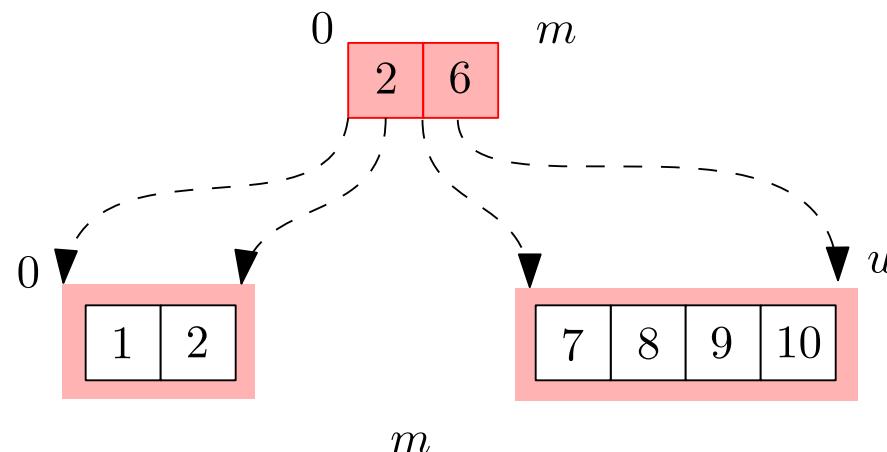
Note that  $\frac{u}{m}$  is the average distance between consecutive elements. It doesn't exploit the distribution of the data, but denser lists require fewer bits.

Some sequence are more compressible than others



# Partitioned Elias-Fano [2]

We can improve compression by exploiting clusters of data with a two-level structure. The first level determines the bounds of the  $b$  clusters, and the second level contains smaller Elias-Fano lists.



## **How can we find the best partitioning to minimize the space occupancy of both levels?**

We can use our partitioning algorithm, assigning a weight to each edge based on the number of bits required to represent the partition in the first level and the Elias-Fano structure in the second level.

The authors also improved the bound by showing that substituting an edge in the path with two sub-edges is always bounded by a constant factor.